

# Naive Bayes Classifier in Python

## Import Libraries

```
In [2]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns

import os
for dirname, _, filenames in os.walk(r"C:\Users\Prachi\Documents\VS Code Files\Machine Learning\Naive Bayes Algorithm\adult.csv"):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

c:\Users\Prachi\anaconda3\Lib\site-packages\pandas\core\arrays\masked.py:61: UserWarning: Pandas requires version '1.3.6' or newer of 'bottleneck' (version '1.3.5' currently installed).

```
from pandas.core import (
C:\Users\Prachi\Documents\VS Code Files\Machine Learning\Naive Bayes Algorithm\adult.csv
C:\Users\Prachi\Documents\VS Code Files\Machine Learning\Naive Bayes Algorithm\adult.csv
ve_bayes_algo.ipynb
```

```
In [3]: import warnings
warnings.filterwarnings('ignore')
```

## Import dataset

```
In [4]: data = 'adult.csv'
df = pd.read_csv(data)
```

```
In [5]: df
```

Out[5]:		age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship
	0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family
	1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family
	2	66	?	186061	Some-college	10	Widowed	?	Unmarried
	3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried
	4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child
	...	...	...	...	...	...	...	...	...
	32556	22	Private	310152	Some-college	10	Never-married	Protective-serv	Not-in-family
	32557	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife
	32558	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband
	32559	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried
	32560	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child

32561 rows × 15 columns



## Exploratory data analysis

In [6]: `df.shape`

Out[6]: (32561, 15)

In [7]: `# preview the dataset`

`df.head()`

Out[7]:

	age	workclass	fnlwgt	education	education.num	marital.status	occupation	relationship	ra
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	Wh
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	Wh
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	Bl
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	Wh
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	Wh

## Rename column names

```
In [8]: col_names = ['age', 'workclass', 'fnlwgt', 'education', 'education_num', 'marital_s',
                    'race', 'sex', 'capital_gain', 'capital_loss', 'hours_per_week', 'nati

df.columns = col_names

df.columns
```

```
Out[8]: Index(['age', 'workclass', 'fnlwgt', 'education', 'education_num',
               'marital_status', 'occupation', 'relationship', 'race', 'sex',
               'capital_gain', 'capital_loss', 'hours_per_week', 'native_country',
               'income'],
              dtype='object')
```

```
In [9]: df.head()
```

Out[9]:

	age	workclass	fnlwgt	education	education_num	marital_status	occupation	relationship	r
0	90	?	77053	HS-grad	9	Widowed	?	Not-in-family	W
1	82	Private	132870	HS-grad	9	Widowed	Exec-managerial	Not-in-family	W
2	66	?	186061	Some-college	10	Widowed	?	Unmarried	B
3	54	Private	140359	7th-8th	4	Divorced	Machine-op-inspct	Unmarried	W
4	41	Private	264663	Some-college	10	Separated	Prof-specialty	Own-child	W

```
In [10]: # view summary of dataset

df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 32561 entries, 0 to 32560
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype
---  -
0   age                   32561 non-null  int64
1   workclass             32561 non-null  object
2   fnlwgt               32561 non-null  int64
3   education             32561 non-null  object
4   education_num        32561 non-null  int64
5   marital_status       32561 non-null  object
6   occupation            32561 non-null  object
7   relationship         32561 non-null  object
8   race                 32561 non-null  object
9   sex                  32561 non-null  object
10  capital_gain         32561 non-null  int64
11  capital_loss         32561 non-null  int64
12  hours_per_week       32561 non-null  int64
13  native_country       32561 non-null  object
14  income               32561 non-null  object
dtypes: int64(6), object(9)
memory usage: 3.7+ MB

```

## Types of variables

### Explore categorical variables

In [11]: *# find categorical variables*

```

categorical = [var for var in df.columns if df[var].dtype=='O']

print('There are {} categorical variables\n'.format(len(categorical)))

print('The categorical variables are :\n\n', categorical)

```

There are 9 categorical variables

The categorical variables are :

```

['workclass', 'education', 'marital_status', 'occupation', 'relationship', 'race', 'sex', 'native_country', 'income']

```

In [12]: *# view the categorical variables*

```

df[categorical].head()

```

Out[12]:	workclass	education	marital_status	occupation	relationship	race	sex	native_country	ii
0		HS-grad	Widowed		Not-in-family	White	Female	United-States	
1	Private	HS-grad	Widowed	Exec-managerial	Not-in-family	White	Female	United-States	
2		Some-college	Widowed		Unmarried	Black	Female	United-States	
3	Private	7th-8th	Divorced	Machine-op-inspct	Unmarried	White	Female	United-States	
4	Private	Some-college	Separated	Prof-specialty	Own-child	White	Female	United-States	

## Explore problems within categorical variables

### Missing values in categorical Variables

In [13]: *# check missing values in categorical variables*

```
df[categorical].isnull().sum()
```

Out[13]:

workclass	0
education	0
marital_status	0
occupation	0
relationship	0
race	0
sex	0
native_country	0
income	0

dtype: int64

### Frequency counts of categorical variables

In [14]: *# view frequency counts of values in categorical variables*

```
for var in categorical:
    print(df[var].value_counts())
```

```

workclass
Private                22696
Self-emp-not-inc      2541
Local-gov             2093
?                     1836
State-gov             1298
Self-emp-inc          1116
Federal-gov           960
Without-pay           14
Never-worked           7
Name: count, dtype: int64
education
HS-grad               10501
Some-college          7291
Bachelors             5355
Masters               1723
Assoc-voc             1382
11th                  1175
Assoc-acdm            1067
10th                  933
7th-8th               646
Prof-school           576
9th                   514
12th                  433
Doctorate             413
5th-6th               333
1st-4th               168
Preschool             51
Name: count, dtype: int64
marital_status
Married-civ-spouse    14976
Never-married         10683
Divorced              4443
Separated             1025
Widowed               993
Married-spouse-absent  418
Married-AF-spouse      23
Name: count, dtype: int64
occupation
Prof-specialty        4140
Craft-repair          4099
Exec-managerial        4066
Adm-clerical          3770
Sales                 3650
Other-service         3295
Machine-op-inspct     2002
?                     1843
Transport-moving      1597
Handlers-cleaners     1370
Farming-fishing        994
Tech-support          928
Protective-serv        649
Priv-house-serv        149
Armed-Forces           9
Name: count, dtype: int64
relationship
Husband               13193
Not-in-family         8305
Own-child             5068
Unmarried             3446
Wife                  1568
Other-relative        981
Name: count, dtype: int64
race

```

```

White                27816
Black                3124
Asian-Pac-Islander   1039
Amer-Indian-Eskimo   311
Other                 271
Name: count, dtype: int64
sex
Male                21790
Female              10771
Name: count, dtype: int64
native_country
United-States       29170
Mexico              643
?                   583
Philippines         198
Germany             137
Canada              121
Puerto-Rico        114
El-Salvador         106
India               100
Cuba                 95
England             90
Jamaica             81
South               80
China               75
Italy               73
Dominican-Republic  70
Vietnam             67
Guatemala           64
Japan               62
Poland              60
Columbia            59
Taiwan              51
Haiti               44
Iran                43
Portugal            37
Nicaragua           34
Peru                31
Greece              29
France              29
Ecuador             28
Ireland             24
Hong                20
Cambodia            19
Trinidad&Tobago     19
Laos                18
Thailand            18
Yugoslavia          16
Outlying-US(Guam-USVI-etc) 14
Hungary             13
Honduras            13
Scotland            12
Holand-Netherlands  1
Name: count, dtype: int64
income
<=50K              24720
>50K                7841
Name: count, dtype: int64

```

In [15]: *# view frequency distribution of categorical variables*

```

for var in categorical:
    print(df[var].value_counts()/float(len(df)))

```

```

workclass
Private          0.697030
Self-emp-not-inc 0.078038
Local-gov        0.064279
?                0.056386
State-gov        0.039864
Self-emp-inc     0.034274
Federal-gov      0.029483
Without-pay      0.000430
Never-worked     0.000215
Name: count, dtype: float64
education
HS-grad          0.322502
Some-college     0.223918
Bachelors        0.164461
Masters          0.052916
Assoc-voc        0.042443
11th             0.036086
Assoc-acdm       0.032769
10th             0.028654
7th-8th          0.019840
Prof-school      0.017690
9th              0.015786
12th             0.013298
Doctorate        0.012684
5th-6th          0.010227
1st-4th          0.005160
Preschool        0.001566
Name: count, dtype: float64
marital_status
Married-civ-spouse 0.459937
Never-married      0.328092
Divorced           0.136452
Separated          0.031479
Widowed            0.030497
Married-spouse-absent 0.012837
Married-AF-spouse  0.000706
Name: count, dtype: float64
occupation
Prof-specialty    0.127146
Craft-repair      0.125887
Exec-managerial   0.124873
Adm-clerical      0.115783
Sales             0.112097
Other-service     0.101195
Machine-op-inspct 0.061485
?                 0.056601
Transport-moving  0.049046
Handlers-cleaners 0.042075
Farming-fishing   0.030527
Tech-support      0.028500
Protective-serv   0.019932
Priv-house-serv   0.004576
Armed-Forces      0.000276
Name: count, dtype: float64
relationship
Husband           0.405178
Not-in-family     0.255060
Own-child         0.155646
Unmarried         0.105832
Wife              0.048156
Other-relative    0.030128
Name: count, dtype: float64
race

```



```

White                0.854274
Black                0.095943
Asian-Pac-Islander  0.031909
Amer-Indian-Eskimo  0.009551
Other                0.008323
Name: count, dtype: float64
sex
Male                0.669205
Female              0.330795
Name: count, dtype: float64
native_country
United-States       0.895857
Mexico              0.019748
?                  0.017905
Philippines         0.006081
Germany             0.004207
Canada              0.003716
Puerto-Rico        0.003501
El-Salvador         0.003255
India               0.003071
Cuba                0.002918
England             0.002764
Jamaica             0.002488
South               0.002457
China               0.002303
Italy               0.002242
Dominican-Republic  0.002150
Vietnam             0.002058
Guatemala           0.001966
Japan               0.001904
Poland              0.001843
Columbia            0.001812
Taiwan              0.001566
Haiti               0.001351
Iran                0.001321
Portugal            0.001136
Nicaragua           0.001044
Peru                0.000952
Greece              0.000891
France              0.000891
Ecuador             0.000860
Ireland             0.000737
Hong                0.000614
Cambodia            0.000584
Trinidad&Tobago     0.000584
Laos                0.000553
Thailand            0.000553
Yugoslavia          0.000491
Outlying-US(Guam-USVI-etc) 0.000430
Hungary             0.000399
Honduras            0.000399
Scotland            0.000369
Holand-Netherlands  0.000031
Name: count, dtype: float64
income
<=50K              0.75919
>50K               0.24081
Name: count, dtype: float64

```

## Explore workclass variable

```
In [16]: # check labels in workclass variable
df.workclass.unique()
```

```
Out[16]: array(['?', 'Private', 'State-gov', 'Federal-gov', 'Self-emp-not-inc',
        'Self-emp-inc', 'Local-gov', 'Without-pay', 'Never-worked'],
        dtype=object)
```

```
In [17]: # check frequency distribution of values in workclass variable

df.workclass.value_counts()
```

```
Out[17]: workclass
Private                22696
Self-emp-not-inc       2541
Local-gov              2093
?                     1836
State-gov              1298
Self-emp-inc           1116
Federal-gov            960
Without-pay            14
Never-worked            7
Name: count, dtype: int64
```

```
In [18]: # replace '?' values in workclass variable with 'NaN'

df['workclass'].replace('?', np.NaN, inplace=True)
```

```
In [19]: # again check the frequency distribution of values in workclass variable

df.workclass.value_counts()
```

```
Out[19]: workclass
Private                22696
Self-emp-not-inc       2541
Local-gov              2093
State-gov              1298
Self-emp-inc           1116
Federal-gov            960
Without-pay            14
Never-worked            7
Name: count, dtype: int64
```

## Explore occupation variable

```
In [20]: # check labels in occupation variable

df.occupation.unique()
```

```
Out[20]: array(['?', 'Exec-managerial', 'Machine-op-inspct', 'Prof-specialty',
        'Other-service', 'Adm-clerical', 'Craft-repair',
        'Transport-moving', 'Handlers-cleaners', 'Sales',
        'Farming-fishing', 'Tech-support', 'Protective-serv',
        'Armed-Forces', 'Priv-house-serv'], dtype=object)
```

```
In [21]: # check frequency distribution of values in occupation variable

df.occupation.value_counts()
```

```
Out[21]: occupation
Prof-specialty      4140
Craft-repair        4099
Exec-managerial     4066
Adm-clerical        3770
Sales               3650
Other-service       3295
Machine-op-inspct   2002
?                  1843
Transport-moving    1597
Handlers-cleaners   1370
Farming-fishing     994
Tech-support        928
Protective-serv     649
Priv-house-serv     149
Armed-Forces        9
Name: count, dtype: int64
```

```
In [22]: # replace '?' values in occupation variable with 'NaN'ArithmeticError

df['occupation'].replace('?', np.NaN, inplace =True)
```

```
In [23]: # again check the frequency distribution of values in occupation variable

df.occupation.value_counts()
```

```
Out[23]: occupation
Prof-specialty      4140
Craft-repair        4099
Exec-managerial     4066
Adm-clerical        3770
Sales               3650
Other-service       3295
Machine-op-inspct   2002
Transport-moving    1597
Handlers-cleaners   1370
Farming-fishing     994
Tech-support        928
Protective-serv     649
Priv-house-serv     149
Armed-Forces        9
Name: count, dtype: int64
```

## Explore native\_country variable

```
In [24]: # check labels in native_country variable

df.native_country.unique()
```

```
Out[24]: array(['United-States', '?', 'Mexico', 'Greece', 'Vietnam', 'China',
        'Taiwan', 'India', 'Philippines', 'Trinidad&Tobago', 'Canada',
        'South', 'Holand-Netherlands', 'Puerto-Rico', 'Poland', 'Iran',
        'England', 'Germany', 'Italy', 'Japan', 'Hong', 'Honduras', 'Cuba',
        'Ireland', 'Cambodia', 'Peru', 'Nicaragua', 'Dominican-Republic',
        'Haiti', 'El-Salvador', 'Hungary', 'Columbia', 'Guatemala',
        'Jamaica', 'Ecuador', 'France', 'Yugoslavia', 'Scotland',
        'Portugal', 'Laos', 'Thailand', 'Outlying-US(Guam-USVI-etc)'],
        dtype=object)
```

```
In [25]: # check frequency distribution of values in native_country variable

df.native_country.value_counts()
```

```

Out[25]: native_country
United-States      29170
Mexico             643
?                  583
Philippines        198
Germany            137
Canada             121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba               95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Guatemala          64
Japan              62
Poland             60
Columbia           59
Taiwan             51
Haiti              44
Iran               43
Portugal           37
Nicaragua          34
Peru               31
Greece            29
France            29
Ecuador           28
Ireland            24
Hong               20
Cambodia           19
Trinidad&Tobago    19
Laos               18
Thailand           18
Yugoslavia         16
Outlying-US(Guam-USVI-etc) 14
Hungary            13
Honduras           13
Scotland           12
Holand-Netherlands 1
Name: count, dtype: int64

```

```

In [26]: # replace '?' values in native_country variable with `NaN`

df['native_country'].replace('?', np.NaN, inplace=True)

```

```

In [27]: # again check the frequency distribution of values in native_country variable

df.native_country.value_counts()

```

```

Out[27]: native_country
United-States      29170
Mexico             643
Philippines        198
Germany            137
Canada             121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba               95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Guatemala          64
Japan              62
Poland             60
Columbia           59
Taiwan             51
Haiti              44
Iran               43
Portugal           37
Nicaragua          34
Peru               31
Greece             29
France             29
Ecuador            28
Ireland            24
Hong               20
Trinidad&Tobago    19
Cambodia           19
Thailand            18
Laos               18
Yugoslavia         16
Outlying-US(Guam-USVI-etc) 14
Hungary            13
Honduras           13
Scotland           12
Holand-Netherlands 1
Name: count, dtype: int64

```

## Check missing values in categorical variables again

```
In [28]: df[categorical].isnull().sum()
```

```

Out[28]: workclass      1836
education      0
marital_status  0
occupation     1843
relationship    0
race           0
sex            0
native_country  583
income         0
dtype: int64

```

# Number of Labels : cardinality

```
In [29]: # check for cardinality in categorical variables

for var in categorical:
    print(var, 'contains', len(df[var].unique()), 'labels')
```

```
workclass contains 9 labels
education contains 16 labels
marital_status contains 7 labels
occupation contains 15 labels
relationship contains 6 labels
race contains 5 labels
sex contains 2 labels
native_country contains 42 labels
income contains 2 labels
```

## Expplre Numerical Variables

```
In [30]: # find numerical variables

numerical = [var for var in df.columns if df[var].dtype!='O']

print('There are {} numerical variables\n'.format(len(numerical)))

print('The numerical variables are :', numerical)
```

There are 6 numerical variables

The numerical variables are : ['age', 'fnlwgt', 'education\_num', 'capital\_gain', 'capital\_loss', 'hours\_per\_week']

```
In [31]: # View the numerical variables
df[numerical].head()
```

```
Out[31]:
```

	age	fnlwgt	education_num	capital_gain	capital_loss	hours_per_week
0	90	77053	9	0	4356	40
1	82	132870	9	0	4356	18
2	66	186061	10	0	4356	40
3	54	140359	4	0	3900	40
4	41	264663	10	0	3900	40

## Explore problems within numerical variables

```
In [32]: # check missing values in numerical variables

df[numerical].isnull().sum()
```

```
Out[32]: age                0
         fnlwgt             0
         education_num      0
         capital_gain        0
         capital_loss        0
         hours_per_week      0
         dtype: int64
```

## Declare feature vector and target variable

```
In [33]: X = df.drop(['income'],axis=1)
         y = df['income']
```

## Split data into separate training and test set

```
In [34]: # split X and y into training and testing sets

         from sklearn.model_selection import train_test_split

         X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_s
```

```
In [35]: # check the shape of X_train and X_test

         X_train.shape, X_test.shape
```

```
Out[35]: ((22792, 14), (9769, 14))
```

## Feature Engineering

```
In [36]: # check data types in X_train

         X_train.dtypes
```

```
Out[36]: age                int64
         workclass          object
         fnlwgt             int64
         education          object
         education_num      int64
         marital_status     object
         occupation         object
         relationship       object
         race               object
         sex                object
         capital_gain        int64
         capital_loss        int64
         hours_per_week      int64
         native_country     object
         dtype: object
```

```
In [37]: # display categorical variables

         categorical = [col for col in X_train.columns if X_train[col].dtypes == 'O']

         categorical
```

```
Out[37]: ['workclass',
          'education',
          'marital_status',
          'occupation',
          'relationship',
          'race',
          'sex',
          'native_country']
```

```
In [38]: # display numerical variables

numerical = [col for col in X_train.columns if X_train[col].dtypes != 'O']

numerical
```

```
Out[38]: ['age',
          'fnlwgt',
          'education_num',
          'capital_gain',
          'capital_loss',
          'hours_per_week']
```

## Engineering missing values in categorical variables

```
In [39]: # print percentage of missing values in the categorical variables in training set

X_train[categorical].isnull().mean()
```

```
Out[39]: workclass      0.056774
          education     0.000000
          marital_status 0.000000
          occupation     0.057038
          relationship    0.000000
          race           0.000000
          sex            0.000000
          native_country  0.018208
          dtype: float64
```

```
In [40]: # print categorical variables with missing data

for col in categorical:
    if X_train[col].isnull().mean()>0:
        print(col, (X_train[col].isnull().mean()))

workclass 0.056774306774306775
occupation 0.057037557037557036
native_country 0.018208143208143207
```

```
In [41]: # impute missing categorical variables with most frequent value

for df2 in [X_train, X_test]:
    df2['workclass'].fillna(X_train['workclass'].mode()[0], inplace=True)
    df2['occupation'].fillna(X_train['occupation'].mode()[0], inplace=True)
    df2['native_country'].fillna(X_train['native_country'].mode()[0], inplace=True)
```

```
In [42]: # check missing values in categorical variables in X_train

X_train[categorical].isnull().sum()
```



```
Out[42]: workclass      0
         education    0
         marital_status 0
         occupation    0
         relationship  0
         race          0
         sex           0
         native_country 0
         dtype: int64
```

```
In [43]: # check missing values in categorical variables in X_test

X_test[categorical].isnull().sum()
```

```
Out[43]: workclass      0
         education    0
         marital_status 0
         occupation    0
         relationship  0
         race          0
         sex           0
         native_country 0
         dtype: int64
```

```
In [44]: # check missing values in X_train

X_train.isnull().sum()
```

```
Out[44]: age          0
         workclass    0
         fnlwgt       0
         education    0
         education_num 0
         marital_status 0
         occupation    0
         relationship  0
         race          0
         sex           0
         capital_gain  0
         capital_loss  0
         hours_per_week 0
         native_country 0
         dtype: int64
```

```
In [45]: # check missing values in X_test

X_test.isnull().sum()
```

```
Out[45]: age          0
         workclass    0
         fnlwgt       0
         education    0
         education_num 0
         marital_status 0
         occupation    0
         relationship  0
         race          0
         sex           0
         capital_gain  0
         capital_loss  0
         hours_per_week 0
         native_country 0
         dtype: int64
```

# Encode categorical Variables

```
In [46]: # print categorical variables
categorical
```

```
Out[46]: ['workclass',
          'education',
          'marital_status',
          'occupation',
          'relationship',
          'race',
          'sex',
          'native_country']
```

```
In [47]: X_train[categorical].head()
```

```
Out[47]:
```

	workclass	education	marital_status	occupation	relationship	race	sex	native_count
32098	State-gov	Bachelors	Married-civ-spouse	Exec-managerial	Wife	White	Female	United-Stat
25206	Local-gov	HS-grad	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	United-Stat
23491	Private	Some-college	Never-married	Exec-managerial	Not-in-family	White	Female	United-Stat
12367	Local-gov	HS-grad	Never-married	Farming-fishing	Own-child	White	Male	United-Stat
7054	Federal-gov	Masters	Married-civ-spouse	Exec-managerial	Husband	White	Male	United-Stat

```
In [48]: # import category encoders
```

```
import category_encoders as ce
```

```
In [49]: # encode remaining variables with one-hot encoding
```

```
encoder = ce.OneHotEncoder(cols=['workclass', 'education', 'marital_status', 'occupation',  
                                'race', 'sex', 'native_country'])
```

```
X_train = encoder.fit_transform(X_train)
```

```
X_test = encoder.transform(X_test)
```

```
In [50]: X_train.head()
```

```
Out[50]:
```

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workcla
<b>32098</b>	40	1	0	0	0	0	0	
<b>25206</b>	39	0	1	0	0	0	0	
<b>23491</b>	42	0	0	1	0	0	0	
<b>12367</b>	27	0	1	0	0	0	0	
<b>7054</b>	38	0	0	0	1	0	0	

5 rows × 105 columns

```
In [51]: X_train.shape
```

```
Out[51]: (22792, 105)
```

```
In [52]: X_test.head()
```

```
Out[52]:
```

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workcla
<b>22278</b>	56	0	0	1	0	0	0	
<b>8950</b>	19	0	0	1	0	0	0	
<b>7838</b>	23	0	0	1	0	0	0	
<b>16505</b>	37	0	0	0	1	0	0	
<b>19140</b>	49	0	0	1	0	0	0	

5 rows × 105 columns

```
In [53]: X_test.shape
```

```
Out[53]: (9769, 105)
```

## Feature Scaling

```
In [54]: cols = X_train.columns
```

```
In [55]: from sklearn.preprocessing import RobustScaler

scaler = RobustScaler()

X_train = scaler.fit_transform(X_train)

X_test = scaler.transform(X_test)
```

```
In [56]: X_train = pd.DataFrame(X_train, columns=[cols])
```

```
In [57]: X_train.head()
```

```
Out[57]:
```

	age	workclass_1	workclass_2	workclass_3	workclass_4	workclass_5	workclass_6	workclass_7
0	0.15	1.0	0.0	-1.0	0.0	0.0	0.0	0.0
1	0.10	0.0	1.0	-1.0	0.0	0.0	0.0	0.0
2	0.25	0.0	0.0	0.0	0.0	0.0	0.0	0.0
3	-0.50	0.0	1.0	-1.0	0.0	0.0	0.0	0.0
4	0.05	0.0	0.0	-1.0	1.0	0.0	0.0	0.0

5 rows × 105 columns

## Model Training

```
In [58]: # train a Gaussian Naive Bayes Classifier on the training set
from sklearn.naive_bayes import GaussianNB

# instantiate the model
gnb = GaussianNB()

# fit the model
gnb.fit(X_train, y_train)
```

```
Out[58]:
```

GaussianNB ⓘ ?

GaussianNB()

## Predict the results

```
In [59]: y_pred = gnb.predict(X_test)
y_pred
```

```
Out[59]: array(['<=50K', '<=50K', '<=50K', ..., '<=50K', '<=50K', '>50K'],
      dtype='<U5')
```

## check the accuracy score

```
In [60]: from sklearn.metrics import accuracy_score

print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred)))

Model accuracy score: 0.8031
```

## compare the train-set and test-set accuracy

```
In [61]: y_pred_train = gnb.predict(X_train)
y_pred_train
```

```
Out[61]: array(['>50K', '<=50K', '<=50K', ..., '<=50K', '>50K', '>50K'],
          dtype='<U5')
```

```
In [62]: print('Training-set accuracy score: {0:0.4f}'.format(accuracy_score(y_train, y_pre
Training-set accuracy score: 0.8009
```

## Check for overfitting and underfitting

```
In [63]: # print the scores on training and test set

print('Training set score : {:.4f}'.format(gnb.score(X_train, y_train)))

print('Test set score : {:.4f}'.format(gnb.score(X_test, y_test)))

Training set score : 0.8009
Test set score : 0.8031
```

## compare model accuracy with null accuracy

```
In [64]: #check class distribution in test set

y_test.value_counts()
```

```
Out[64]: income
<=50K    7410
>50K     2359
Name: count, dtype: int64
```

```
In [65]: # check null accuracy score

null_accuracy = (7410/(7410+2359))

print('Null accuracy score: {0:0.4f}'.format(null_accuracy))

Null accuracy score: 0.7585
```

## confusion matrix

```
In [66]: # Print the Confusion Matrix and slice it into four pieces

from sklearn.metrics import confusion_matrix

cm = confusion_matrix(y_test, y_pred)

print('Confusion matrix\n\n', cm)

print('\nTrue Positives(TP) = ', cm[0,0])

print('\nTrue Negatives(TN) = ', cm[1,1])

print('\nFalse Positives(FP) = ', cm[0,1])

print('\nFalse Negatives(FN) = ', cm[1,0])
```

Confusion matrix

```
[[5953 1457]
 [ 467 1892]]
```

True Positives(TP) = 5953

True Negatives(TN) = 1892

False Positives(FP) = 1457


False Negatives(FN) = 467

```
In [67]: # visualize confusion matrix with seaborn heatmap

cm_matrix = pd.DataFrame(data=cm, columns=['Actual Positive:1', 'Actual Negative:0'],
                          index=['Predict Positive:1', 'Predict Negative:0'])

sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

Out[67]: <Axes: >



## Classification metrics

## Classification Report

```
In [68]: from sklearn.metrics import classification_report

print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
<=50K	0.93	0.80	0.86	7410
>50K	0.56	0.80	0.66	2359
accuracy			0.80	9769
macro avg	0.75	0.80	0.76	9769
weighted avg	0.84	0.80	0.81	9769

## Classification Accuracy

```
In [69]: TP = cm[0,0]
         TN = cm[1,1]
         FP = cm[0,1]
         FN = cm[1,0]
```

## Classification error

```
In [70]: # print classification error

classification_error = (FP + FN) / float(TP + TN + FP + FN)

print('Classification error : {0:0.4f}'.format(classification_error))

Classification error : 0.1969
```

## Precision

```
In [71]: # print precision score

precision = TP / float(TP+FP)

print('Precision: {0:0.4f}'.format(precision))

Precision: 0.8034
```

## Recall

```
In [73]: recall = TP / float(TP+FN)

print('Recall or Sensitivity : {0:0.4f}'.format(recall))

Recall or Sensitivity : 0.9273
```

## True Positive Rate

```
In [74]: true_positive_rate = TP / float(TP+FN)

print('True Positive Rate: {0:0.4f}'.format(true_positive_rate))

True Positive Rate: 0.9273
```

# False Positive Rate

```
In [75]: false_positive_rate = FP/float(FP+TN)

print('False Positive Rate: {0:0.4f}'.format(false_positive_rate))

False Positive Rate: 0.4351
```

# Specificity

```
In [ ]: specificity = TN / (TN + FP)

print('Specificity : {0:0.4f}'.format(specificity))
```

# F1 -Score

# Calculate class probabilities

```
In [76]: # print the first 10 predicted probabilities of two classes- 0 and 1

y_pred_prob = gnb.predict_proba(X_test)[0:10]

y_pred_prob
```

```
Out[76]: array([[9.99999693e-01, 3.06618197e-07],
 [1.00000000e+00, 1.02355439e-10],
 [9.99999997e-01, 3.02850706e-09],
 [8.78002299e-04, 9.99121998e-01],
 [7.55021219e-04, 9.99244979e-01],
 [9.99505992e-01, 4.94008099e-04],
 [9.99999697e-01, 3.03376335e-07],
 [9.63760637e-01, 3.62393626e-02],
 [9.99999937e-01, 6.31028512e-08],
 [1.41650243e-03, 9.98583498e-01]])
```

```
In [77]: # store the probabilities in dataframe

y_pred_prob_df = pd.DataFrame(data=y_pred_prob, columns=['Prob of - <=50K', 'Prob of >50K'])

y_pred_prob_df
```



Out[77]:

	Prob of - <=50K	Prob of - >50K
0	1.000000	3.066182e-07
1	1.000000	1.023554e-10
2	1.000000	3.028507e-09
3	0.000878	9.991220e-01
4	0.000755	9.992450e-01
5	0.999506	4.940081e-04
6	1.000000	3.033763e-07
7	0.963761	3.623936e-02
8	1.000000	6.310285e-08
9	0.001417	9.985835e-01

In [78]: *# print the first 10 predicted probabilities for class 1 - Probability of >50K*

```
gnb.predict_proba(X_test)[0:10, 1]
```

Out[78]: array([3.06618197e-07, 1.02355439e-10, 3.02850706e-09, 9.99121998e-01,  
9.99244979e-01, 4.94008099e-04, 3.03376335e-07, 3.62393626e-02,  
6.31028512e-08, 9.98583498e-01])

In [79]: *# store the predicted probabilities for class 1 - Probability of >50K*

```
y_pred1 = gnb.predict_proba(X_test)[: , 1]
```

In [80]: *# plot histogram of predicted probabilities*

```
# adjust the font size
```

```
plt.rcParams['font.size'] = 12
```

```
# plot histogram with 10 bins
```

```
plt.hist(y_pred1, bins = 10)
```

```
# set the title of predicted probabilities
```

```
plt.title('Histogram of predicted probabilities of salaries >50K')
```

```
# set the x-axis limit
```

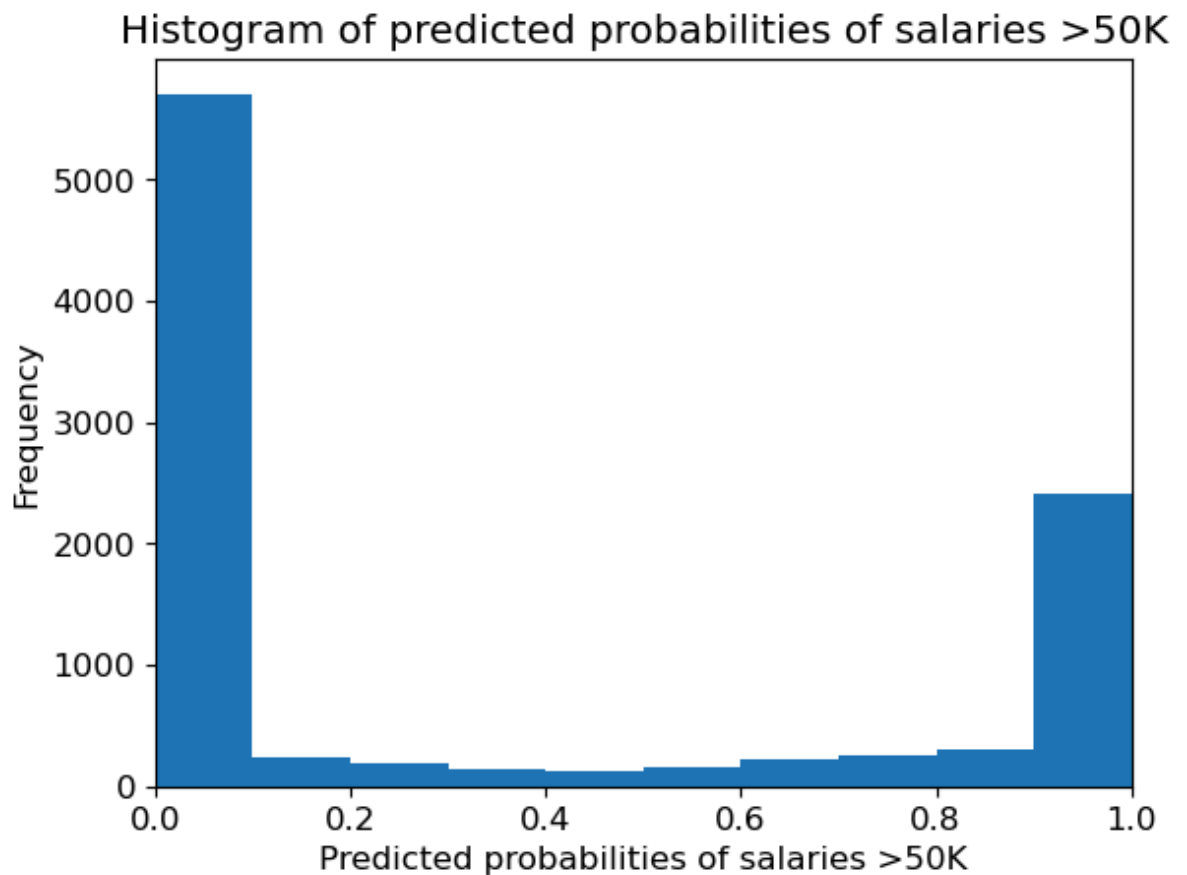
```
plt.xlim(0,1)
```

```
# set the title
```

```
plt.xlabel('Predicted probabilities of salaries >50K')
```

```
plt.ylabel('Frequency')
```

Out[80]: Text(0, 0.5, 'Frequency')



## ROC AUC

```
In [81]: # plot ROC Curve

from sklearn.metrics import roc_curve

fpr, tpr, thresholds = roc_curve(y_test, y_pred1, pos_label = '>50K')

plt.figure(figsize=(6,4))

plt.plot(fpr, tpr, linewidth=2)

plt.plot([0,1], [0,1], 'k--' )

plt.rcParams['font.size'] = 12


plt.title('ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries')

plt.xlabel('False Positive Rate (1 - Specificity)')

plt.ylabel('True Positive Rate (Sensitivity)')

plt.show()
```

## ROC curve for Gaussian Naive Bayes Classifier for Predicting Salaries



```
In [82]: # compute ROC AUC

from sklearn.metrics import roc_auc_score

ROC_AUC = roc_auc_score(y_test, y_pred1)

print('ROC AUC : {:.4f}'.format(ROC_AUC))

ROC AUC : 0.8909
```

```
In [83]: # calculate cross-validated ROC AUC

from sklearn.model_selection import cross_val_score

Cross_validated_ROC_AUC = cross_val_score(gnb, X_train, y_train, cv=5, scoring='roc_auc')

print('Cross validated ROC AUC : {:.4f}'.format(Cross_validated_ROC_AUC))

Cross validated ROC AUC : 0.8936
```

## K-fold Cross validation

```
In [84]: # Applying 10-Fold Cross Validation

from sklearn.model_selection import cross_val_score

scores = cross_val_score(gnb, X_train, y_train, cv = 10, scoring='accuracy')

print('Cross-validation scores:{}'.format(scores))

Cross-validation scores:[0.80701754 0.7877193  0.79947345 0.81439228 0.785871  0.81526986
 0.78894252 0.79420799 0.80122861 0.8056165 ]
```

```
In [85]: # compute Average cross-validation score

print('Average cross-validation score: {:.4f}'.format(scores.mean()))

Average cross-validation score: 0.8000
```

# Results and Conclusion

1. In this project, I build a Gaussian Naïve Bayes Classifier model to predict whether a person makes over 50K a year. The model yields a very good performance as indicated by the model accuracy which was found to be 0.8083.
2. The training-set accuracy score is 0.8067 while the test-set accuracy to be 0.8083. These two values are quite comparable. So, there is no sign of overfitting.
3. I have compared the model accuracy score which is 0.8083 with null accuracy score which is 0.7582. So, we can conclude that our Gaussian Naïve Bayes classifier model is doing a very good job in predicting the class labels.
4. ROC AUC of our model approaches towards 1. So, we can conclude that our classifier does a very good job in predicting whether a person makes over 50K a year.
5. Using the mean cross-validation, we can conclude that we expect the model to be around 80.63% accurate on average.
6. If we look at all the 10 scores produced by the 10-fold cross-validation, we can also conclude that there is a relatively small variance in the accuracy between folds, ranging from 81.35% accuracy to 79.64% accuracy. So, we can conclude that the model is independent of the particular folds used for training.
7. Our original model accuracy is 0.8083, but the mean cross-validation accuracy is 0.8063. So, the 10-fold cross-validation accuracy does not result in performance improvement for this model.