

```
In [5]: import os
os.getcwd()
```

```
Out[5]: 'c:\\Users\\Prachi\\Documents\\VS Code Files\\ML CAPSTONE PROJECT\\Diabetes Prediction'
```

```
In [6]: os.chdir(r"C:\Users\Prachi\Documents\VS Code Files\ML CAPSTONE PROJECT\Diabetes Pre
```

```
In [7]: import numpy as np

import pandas as pd #excellent for dataset manipulation

# for data visulization
import matplotlib.pyplot as plt

#stats visualization
import seaborn as sns

#Labelencoding to convert categorical data into lowlevel language
from sklearn.preprocessing import LabelEncoder

#scaling data
from sklearn.preprocessing import StandardScaler

#data partions
from sklearn.model_selection import train_test_split

#algorithams
from sklearn.linear_model import LogisticRegression

from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import RandomForestClassifier

from xgboost import XGBClassifier

#accuracy confusion matric and classification report
from sklearn.metrics import accuracy_score, confusion_matrix, classification_report

import warnings

# To ignore all warnings
warnings.filterwarnings("ignore")
```

```
In [8]: df = pd.read_csv("diabetes_prediction_dataset.csv") # Reading csv file
```

```
In [9]: df.head()
```

```
Out[9]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_level
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	No Info	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	

```
In [10]: df.isna().any() # checking is there any null values
```

```
Out[10]:
```

gender	False
age	False
hypertension	False
heart_disease	False
smoking_history	False
bmi	False
HbA1c_level	False
blood_glucose_level	False
diabetes	False
dtype:	bool

```
In [11]: df.corr(numeric_only=True) # correlation
```

```
Out[11]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
age	1.000000	0.251171	0.233354	0.337396	0.101354	0.110
hypertension	0.251171	1.000000	0.121262	0.147666	0.080939	0.084
heart_disease	0.233354	0.121262	1.000000	0.061198	0.067589	0.070
bmi	0.337396	0.147666	0.061198	1.000000	0.082997	0.091
HbA1c_level	0.101354	0.080939	0.067589	0.082997	1.000000	0.166
blood_glucose_level	0.110672	0.084429	0.070066	0.091261	0.166733	1.000
diabetes	0.258008	0.197823	0.171727	0.214357	0.400660	0.419

```
In [12]: df.shape # shape of the dataframe
```

```
Out[12]: (100000, 9)
```

```
In [13]: for column in df.columns:
          unique_values = df[column].unique()

          # printing unique values
          print('Column "{}" has unique values: {}'.format(column, unique_values))
```

Column "gender" has unique values: ['Female' 'Male' 'Other']  
 Column "age" has unique values: [80. 54. 28. 36. 76. 20. 44. 79. 4  
 2. 32. 53. 78.  
 67. 15. 37. 40. 5. 69. 72. 4. 30. 45. 43. 50.  
 41. 26. 34. 73. 77. 66. 29. 60. 38. 3. 57. 74.  
 19. 46. 21. 59. 27. 13. 56. 2. 7. 11. 6. 55.  
 9. 62. 47. 12. 68. 75. 22. 58. 18. 24. 17. 25.  
 0.08 33. 16. 61. 31. 8. 49. 39. 65. 14. 70. 0.56  
 48. 51. 71. 0.88 64. 63. 52. 0.16 10. 35. 23. 0.64  
 1.16 1.64 0.72 1.88 1.32 0.8 1.24 1. 1.8 0.48 1.56 1.08  
 0.24 1.4 0.4 0.32 1.72 1.48]  
 Column "hypertension" has unique values: [0 1]  
 Column "heart\_disease" has unique values: [1 0]  
 Column "smoking\_history" has unique values: ['never' 'No Info' 'current' 'former'  
 'ever' 'not current']  
 Column "bmi" has unique values: [25.19 27.32 23.45 ... 59.42 44.39 60.52]  
 Column "HbA1c\_level" has unique values: [6.6 5.7 5. 4.8 6.5 6.1 6. 5.8 3.5 6.2  
 4. 4.5 9. 7. 8.8 8.2 7.5 6.8]  
 Column "blood\_glucose\_level" has unique values: [140 80 158 155 85 200 145 100 1  
 30 160 126 159 90 260 220 300 280 240]  
 Column "diabetes" has unique values: [0 1]

```
In [14]: df['smoking_history'].value_counts() # value count of smoking_history parameter i
```

```
Out[14]: smoking_history
No Info      35816
never        35095
former        9352
current      9286
not current  6447
ever         4004
Name: count, dtype: int64
```

```
In [15]: df["smoking_history"].value_counts()/len(df) #finding the percentage
```

```
Out[15]: smoking_history
No Info      0.35816
never        0.35095
former        0.09352
current      0.09286
not current  0.06447
ever         0.04004
Name: count, dtype: float64
```

```
In [16]: # Replaceing No Info columns with pd.NA
df['smoking_history'] = df['smoking_history'].replace('No Info', pd.NA)

# Replace missing values with the mode it is string so we are using mode
mode_value = df['smoking_history'].mode()[0]
df['smoking_history'] = df['smoking_history'].fillna(mode_value) #filling no info v

# Printing the updated value counts
print(df['smoking_history'].value_counts())
```

```
smoking_history
never        70911
former        9352
current      9286
not current  6447
ever         4004
Name: count, dtype: int64
```

```
In [17]: df.info() # information of the dataframe
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   gender                 100000 non-null object
1   age                   100000 non-null float64
2   hypertension          100000 non-null int64
3   heart_disease         100000 non-null int64
4   smoking_history       100000 non-null object
5   bmi                   100000 non-null float64
6   HbA1c_level           100000 non-null float64
7   blood_glucose_level   100000 non-null int64
8   diabetes              100000 non-null int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB

```

```
In [18]: df.gender.value_counts()
```

```

Out[18]: gender
Female    58552
Male      41430
Other       18
Name: count, dtype: int64

```

```
In [19]: df.describe()
```

```

Out[19]:

```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
count	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000	100000.000000
mean	41.885856	0.07485	0.039420	27.320767	5.527507	138.058
std	22.516840	0.26315	0.194593	6.636783	1.070672	40.708
min	0.080000	0.00000	0.000000	10.010000	3.500000	80.000
25%	24.000000	0.00000	0.000000	23.630000	4.800000	100.000
50%	43.000000	0.00000	0.000000	27.320000	5.800000	140.000
75%	60.000000	0.00000	0.000000	29.580000	6.200000	159.000
max	80.000000	1.00000	1.000000	95.690000	9.000000	300.000

```

In [20]: #removing , in bmi parameter
df["bmi"] = [float(str(i).replace(",","")) for i in df["bmi"]]

```

```
In [21]: df
```

Out[21]:

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose
0	Female	80.0	0	1	never	25.19	6.6	
1	Female	54.0	0	0	never	27.32	6.6	
2	Male	28.0	0	0	never	27.32	5.7	
3	Female	36.0	0	0	current	23.45	5.0	
4	Male	76.0	1	1	current	20.14	4.8	
...	...	...	...	...	...	...	...	...
99995	Female	80.0	0	0	never	27.32	6.2	
99996	Female	2.0	0	0	never	17.37	6.5	
99997	Male	66.0	0	0	former	27.83	5.7	
99998	Female	24.0	0	0	never	35.42	4.0	
99999	Female	57.0	0	0	current	22.43	6.6	

100000 rows × 9 columns

```
In [22]: # plotting value_counts of diabetes in graphical representation
df['diabetes'].value_counts().plot(kind='barh')


#Xlabel name
plt.xlabel('count')

#ylabel name
plt.ylabel('diabetes')

#title of the plot
plt.title('count of diabetes and Non diabetes')

#invert ylabes to no diabetes on top
plt.gca().invert_yaxis()

#printing the plot
plt.show()
```



```
In [23]: df['diabetes'].value_counts()/len(df) #percentage of 1--diabetes and 2--no diabetes
```

```
Out[23]: diabetes
0    0.915
1    0.085
Name: count, dtype: float64
```

```
In [24]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   gender                 100000 non-null  object
1   age                   100000 non-null  float64
2   hypertension           100000 non-null  int64
3   heart_disease          100000 non-null  int64
4   smoking_history        100000 non-null  object
5   bmi                   100000 non-null  float64
6   HbA1c_level            100000 non-null  float64
7   blood_glucose_level    100000 non-null  int64
8   diabetes               100000 non-null  int64
dtypes: float64(3), int64(4), object(2)
memory usage: 6.9+ MB
```

```
In [25]: le = LabelEncoder() # activating label encoder function
le
```

```
Out[25]: ▼ LabelEncoder ⓘ ?
          ► Parameters
```

```
In [26]: Label_encod_columns=['gender','smoking_history'] #selecting columns to apply label
df[Label_encod_columns]=df[Label_encod_columns].apply(le.fit_transform) #applying l
```


```
In [27]: df.head(3)
```

```
Out[27]:
```

	gender	age	hypertension	heart_disease	smoking_history	bmi	HbA1c_level	blood_glucose_
0	0	80.0	0	1	3	25.19	6.6	
1	0	54.0	0	0	3	27.32	6.6	
2	1	28.0	0	0	3	27.32	5.7	


```
In [28]: sns.boxplot(data=df[['age','blood_glucose_level','bmi']]) #checking outliers using
```

```
Out[28]: <Axes: >
```



```
In [29]: sns.boxplot(data=df['HbA1c_level']) #checking outlayers using boxplot
```

```
Out[29]: <Axes: >
```




```
In [30]: ''' it is always good to ignore outliers in medical data '''
```

```
Out[30]: ' it is always good to ignore outliers in medical data '
```


```
In [31]: sns.lmplot(data=df, x='blood_glucose_level', y='diabetes', fit_reg=False)#lmplot plot
```

```
Out[31]: <seaborn.axisgrid.FacetGrid at 0x1fc3045f350>
```





```
In [32]: sns.pairplot(df) #using pairplot to check relation between parameters  
  
         #print the pairplot  
         plt.show()
```




In [33]: `df.corr()`

Out[33]:

	gender	age	hypertension	heart_disease	smoking_history	bmi
gender	1.000000	-0.030656	0.014203	0.077696	-0.044081	-0.022994
age	-0.030656	1.000000	0.251171	0.233354	-0.098969	0.337396
hypertension	0.014203	0.251171	1.000000	0.121262	-0.048631	0.147666
heart_disease	0.077696	0.233354	0.121262	1.000000	-0.048253	0.061198
smoking_history	-0.044081	-0.098969	-0.048631	-0.048253	1.000000	-0.087735
bmi	-0.022994	0.337396	0.147666	0.061198	-0.087735	1.000000
HbA1c_level	0.019957	0.101354	0.080939	0.067589	-0.017534	0.082997
blood_glucose_level	0.017199	0.110672	0.084429	0.070066	-0.022985	0.091261
diabetes	0.037411	0.258008	0.197823	0.171727	-0.049841	0.214357

In [34]: `plt.figure(figsize=(20,8)) #figsize`  
*#printing graphical representations of*  
`df.corr()['diabetes'].sort_values(ascending=False).plot(kind='bar')`

Out[34]: <Axes: >



In [35]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100000 entries, 0 to 99999
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   gender                100000 non-null int32   
1   age                  100000 non-null float64  
2   hypertension         100000 non-null int64   
3   heart_disease        100000 non-null int64   
4   smoking_history      100000 non-null int32   
5   bmi                  100000 non-null float64  
6   HbA1c_level          100000 non-null float64  
7   blood_glucose_level  100000 non-null int64   
8   diabetes             100000 non-null int64   
dtypes: float64(3), int32(2), int64(4)
memory usage: 6.1 MB
```

In [36]: `# selecting x variables`

```
#selecting X variables
X = df.loc[:, 'age':'heart_disease'].join(df.loc[:, 'bmi':'blood_glucose_level'])
```

X

```
Out[36]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
0	80.0	0	1	25.19	6.6	140
1	54.0	0	0	27.32	6.6	80
2	28.0	0	0	27.32	5.7	158
3	36.0	0	0	23.45	5.0	155
4	76.0	1	1	20.14	4.8	155
...	...	...	...	...	...	...
99995	80.0	0	0	27.32	6.2	90
99996	2.0	0	0	17.37	6.5	100
99997	66.0	0	0	27.83	5.7	155
99998	24.0	0	0	35.42	4.0	100
99999	57.0	0	0	22.43	6.6	90

100000 rows × 6 columns

```
In [37]: y=df.loc[:, 'diabetes'] #y variable
y #printing y variable
```

```
Out[37]:
```

0	0
1	0
2	0
3	0
4	0
...	..
99995	0
99996	0
99997	0
99998	0
99999	0

Name: diabetes, Length: 100000, dtype: int64

```
In [38]: # splitting trining and testing data in 70 30 rating testing size is 0.3 random_stat
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.2,random_state=0)
```

```
In [39]: X_train.head()
```

```
Out[39]:
```

	age	hypertension	heart_disease	bmi	HbA1c_level	blood_glucose_level
10382	2.0	0	0	16.45	6.2	159
73171	55.0	0	0	24.59	6.0	130
30938	24.0	0	0	21.77	4.5	130
99310	30.0	0	0	27.32	6.2	159
58959	13.0	0	0	18.37	6.5	130

```
In [40]: print('Shape of Train data')
print(X_train.shape)
```

```

print(y_train.shape)

print('Shape of Testing data')

print(X_test.shape)

print(y_test.shape)

```

```

Shape of Train data
(80000, 6)
(80000,)
Shape of Testing data
(20000, 6)
(20000,)

```

```

In [41]: ss = StandardScaler() # activating StandardScaler
ss

```

```

Out[41]:
StandardScaler ⓘ ⓘ
Parameters

```

```

In [42]: X_train_scaled=ss.fit_transform(X_train) #scaling X_train data

```

```

In [43]: if len(X_test.shape) == 1: #if x is 1d array
          X_test = X_test.values.reshape(-1, 1) #converting to 2d array

          X_test_scaled = ss.fit_transform(X_test) #scaling X_test data

```

```

In [44]: model_lr = LogisticRegression() # activating Logistic regression

```

```

In [45]: model_lr.fit(X_train_scaled,y_train) # training logistic regression model

```

```

Out[45]:
LogisticRegression ⓘ ⓘ
Parameters

```

```

In [46]: y_pred = model_lr.predict(X_test_scaled) #predecting y_test data
          y_pred[:10]

```

```

Out[46]: array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0], dtype=int64)

```

```

In [47]: y_test[:10] # actual y_test data

```

```

Out[47]:
3582      0
60498     0
53227     0
21333     0
3885      0
51521     0
84261     0
10685     1
59948     0
41032     0
Name: diabetes, dtype: int64

```

```

In [48]: accuracy_score(y_pred,y_test) #accuracy_score

```

Out[48]: 0.95975

```
In [49]: print(classification_report(y_pred,y_test))    # classification report
```

	precision	recall	f1-score	support
0	0.99	0.97	0.98	18736
1	0.63	0.86	0.73	1264
accuracy			0.96	20000
macro avg	0.81	0.91	0.85	20000
weighted avg	0.97	0.96	0.96	20000

```
In [50]: confusion_matrix(y_pred,y_test) #confusion_matrix
```

```
Out[50]: array([[18114,  622],
               [ 183, 1081]], dtype=int64)
```

```
In [51]: y_train.value_counts()    # data is highly imbalancing
```

```
Out[51]: diabetes
0      73203
1       6797
Name: count, dtype: int64
```

```
In [52]: value_counts = y_train.value_counts()


plt.figure(figsize=(16,8))

plt.pie(value_counts, labels=value_counts.index, autopct='%1.2f%%', startangle=140)

plt.title('Distribution of y_train')

plt.axis('equal')    # Equal aspect ratio ensures that pie is drawn as a circle

plt.show()
```



```
In [53]: from imblearn.over_sampling import SMOTE # using smote function to balance our set

smote=SMOTE()

X_ovs,y_ovs=smote.fit_resample(X,y) #passing X and y variables to it to balance out
```


```
fig, oversp = plt.subplots()

oversp.pie( y_ovs.value_counts(), autopct='%.2f')

oversp.set_title("Over-sampling")

plt.show()
```

Over-sampling



```
In [54]: # Dividing our resampling data into 70 30 ratio

Xr_train,Xr_test,yr_train,yr_test=train_test_split(X_ovs,y_ovs,train_size=0.7,randc
```

```
In [55]: print('train data shape')

print(Xr_train.shape)

print(yr_train.shape)

print('test data shape')

print(Xr_test.shape)

print(yr_test.shape)
```

```
train data shape
(128099, 6)
(128099,)
test data shape
(54901, 6)
(54901,)
```

```
In [56]: print('y_train and y_test value_count')
print(yr_train.value_counts())
print(yr_test.value_counts())
```

```
y_train and y_test value_count
diabetes
0    64131
1    63968
Name: count, dtype: int64
diabetes
1    27532
0    27369
Name: count, dtype: int64
```

```
In [57]: ss = StandardScaler()
ss
```

```
Out[57]: ▼ StandardScaler ⓘ ?
► Parameters
```

```
In [59]: data = Xr_train, Xr_test

xr_train_sc = ss.fit_transform(Xr_train) # scaling oyr resampling data xr trsin
Xr_test_sc = ss.fit_transform(Xr_test) # scaling our sampling Xr_test data
```

```
In [60]: Xr_train_scaled = pd.DataFrame(xr_train_sc) #Xr_train_scaled converting into the da

print(Xr_train_scaled.shape)
Xr_train_scaled.head()
print(yr_train.shape)

(128099, 6)
(128099,)
```

```
In [61]: Xr_test_scaled=pd.DataFrame(Xr_test_sc) #Xr_test converting into the dataframe

print(Xr_test_scaled.shape)
Xr_test_scaled.head()

(54901, 6)
```

```
Out[61]:
```

	0	1	2	3	4	5
0	-0.190968	-0.294845	-0.204944	-0.169632	1.914409	2.037921
1	-1.098037	-0.294845	-0.204944	-0.404830	0.373911	-0.060991
2	-1.469745	-0.294845	-0.204944	-0.287753	0.373911	-1.460265
3	-0.772793	3.391618	-0.204944	0.289464	0.373911	-1.372811
4	-1.376818	-0.294845	-0.204944	-0.287753	-2.153574	-1.110447

```
In [63]: model_lk = LogisticRegression()

model_lk.fit(Xr_train_scaled, yr_train)
```

```
Out[63]: ▼ LogisticRegression ⓘ ?
► Parameters
```

```
In [64]: y_pred_lr=model_lk.predict(Xr_test_scaled) #predecting yr_test data
y_pred_lr[:10]
```



```
Out[64]: array([1, 0, 0, 0, 0, 1, 0, 0, 0, 1], dtype=int64)
```

```
In [65]: yr_test[:10]
```

```
Out[65]: 180328    1
          573      0
          13494    0
          93981    0
          75389    0
          180973    1
          71021    0
          19293    0
          16393    0
          121419    1
          Name: diabetes, dtype: int64
```

```
In [66]: #classification_report for predict value and orginal value

print(classification_report(y_pred_lr, yr_test))
```

	precision	recall	f1-score	support
0	0.88	0.88	0.88	27321
1	0.89	0.88	0.88	27580
accuracy			0.88	54901
macro avg	0.88	0.88	0.88	54901
weighted avg	0.88	0.88	0.88	54901

```
In [67]: #confusion_matrix for predict value and orginal value

confusion_matrix(y_pred_lr, yr_test)
```

```
Out[67]: array([[24155,  3166],
                [ 3214, 24366]], dtype=int64)
```

## Decision Tree Classifier

```
In [68]: # activating DecisionTree Classifier
model_dtc=DecisionTreeClassifier()

# passing xr_train_scaled, yr_train to training the model
model_dtc.fit(Xr_train_scaled, yr_train)

model_dtc
```

```
Out[68]: ▾ DecisionTreeClassifier ⓘ ?
          ► Parameters
```

```
In [69]: y_pred_dtc = model_dtc.predict(Xr_test_scaled)    # predicting yr_test data
```

```
In [70]: # classification report for decisionTreeClassifier

print(classification_report(y_pred_dtc, yr_test))
```

	precision	recall	f1-score	support
0	0.63	1.00	0.77	17293
1	1.00	0.73	0.84	37608
accuracy			0.81	54901
macro avg	0.81	0.86	0.81	54901
weighted avg	0.88	0.81	0.82	54901

```
In [71]: confusion_matrix(y_pred_dtc,yr_test)
```

```
Out[71]: array([[17229,    64],
               [10140, 27468]], dtype=int64)
```

## RandomForest Classifier

```
In [73]: model_rfc = RandomForestClassifier()

model_rfc.fit(Xr_train_scaled, yr_train)
```

```
Out[73]: ▼ RandomForestClassifier ⓘ ?
```

► Parameters

```
In [74]: y_pred_rfc = model_rfc.predict(Xr_test_scaled)
```

```
In [75]: print(classification_report(y_pred_rfc,yr_test))
```

	precision	recall	f1-score	support
0	0.77	0.99	0.87	21302
1	0.99	0.81	0.89	33599
accuracy			0.88	54901
macro avg	0.88	0.90	0.88	54901
weighted avg	0.91	0.88	0.88	54901

```
In [76]: confusion_matrix(y_pred_rfc,yr_test)
```

```
Out[76]: array([[21123,   179],
               [ 6246, 27353]], dtype=int64)
```

## XGBoost

```
In [77]: model_xgb=XGBClassifier()

model_xgb.fit(Xr_train_scaled,yr_train)
```

```
Out[77]: ▼ XGBClassifier ⓘ ?
```

► Parameters

```
In [78]: y_pred_xgb = model_xgb.predict(Xr_test_scaled)
```

```
In [79]: print(classification_report(y_pred_xgb,yr_test))
```

	precision	recall	f1-score	support
0	0.90	0.96	0.93	25811
1	0.96	0.91	0.93	29090
accuracy			0.93	54901
macro avg	0.93	0.93	0.93	54901
weighted avg	0.93	0.93	0.93	54901

```
In [80]: confusion_matrix(y_pred_xgb,yr_test)
```

```
Out[80]: array([[24737, 1074],  
               [ 2632, 26458]], dtype=int64)
```

```
In [81]: from sklearn.model_selection import GridSearchCV, cross_val_score  
from sklearn.linear_model import LogisticRegression  
  
# Define the parameter grid to search over  
param_grid = {  
    'C': [0.001, 0.01, 0.1, 1, 10, 100], # Regularization parameter  
    'penalty': ['l1', 'l2'] # Penalty type  
}  
  
# Create a Logistic Regression model  
logistic = LogisticRegression()  
  
# Create a GridSearchCV object  
grid_search = GridSearchCV(estimator=logistic, param_grid=param_grid, cv=10)  
  
# Initialize an empty list to store the accuracy scores  
accuracy_scores = []  
  
# Perform cross-validation 10 times  
for _ in range(10):  
    # Fit the GridSearchCV object to the training data  
    grid_search.fit(Xr_train_scaled, yr_train)  
  
    # Get the best parameters  
    best_params = grid_search.best_params_  
  
    # Perform cross-validation with the best model  
    cv_scores = cross_val_score(grid_search.best_estimator_, Xr_train_scaled, yr_train,  
                                scoring='accuracy', cv=5)  
  
    # Store the mean accuracy score  
    accuracy_scores.append(cv_scores.mean())  
  
# Print the accuracy scores obtained over 10 iterations  
# print("Accuracy scores over 10 iterations:", accuracy_scores)  
print("Accuracy scores over 10 iterations:", ["{:.2f}".format(score) for score in accuracy_scores])  
  
# Get the best parameters and best score  
best_params = grid_search.best_params_  
best_score = grid_search.best_score_  
  
print("Best parameters found:", best_params)  
print("Best cross-validation score:", best_score)
```

Accuracy scores over 10 iterations: ['0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89', '0.89']

Best parameters found: {'C': 0.001, 'penalty': 'l2'}

Best cross-validation score: 0.8851200978478675

## Final Model

```
In [82]: from sklearn.linear_model import LogisticRegression

# Create a Logistic Regression model with the best parameters
final_model = LogisticRegression(C=0.001, penalty='l2')

# Fit the final model to the entire training dataset
final_model.fit(Xr_train_scaled, yr_train)
```

```
Out[82]: ▾ LogisticRegression ⓘ ?
          ► Parameters
```

```
In [83]: import pickle

# Save the final model to a pickle file
with open('final_model.pkl', 'wb') as file:
    pickle.dump(final_model, file)
```

```
In [84]: import pickle
import numpy as np

# Load the model from the pickle file
with open('final_model.pkl', 'rb') as file:
    loaded_model = pickle.load(file)

# Define the mean and standard deviation of the training data
mean_values = [41.885856, 0.07485, 0.03942, 27.320767, 5.527507, 138.058060]
std_values = [22.516840, 0.26315, 0.194593, 6.636783, 1.070672, 40.708136]

# Define the input features for prediction
age = 30
hypertension = 0
heart_disease = 0
bmi = 100.0
HbA1c_level = 5.0
blood_glucose_level = 90

# Scale the input features manually
scaled_features = [(x - mean) / std for x, mean, std in zip(
    [age, hypertension, heart_disease, bmi, HbA1c_level, blood_glucose_level],
    mean_values, std_values
)]

# Make predictions on the scaled data
prediction = loaded_model.predict([scaled_features])

# Print the prediction
if prediction[0] == 1:
    print("Diabetic")
else:
    print("Not Diabetic")
```

Diabetic